

# FinanceBrain - Memory-Persistent Research Assistant

## 1. Introduction

This document outlines the comprehensive architecture and design decisions for **FinanceBrain**, a memory-persistent research assistant built using LlamaIndex Workflows. The system is designed to process complex user queries about financial documents (like **Adobe Annual Report**), maintain conversational context across sessions, and provide intelligent, contextually-aware responses through advanced query decomposition and content analysis.

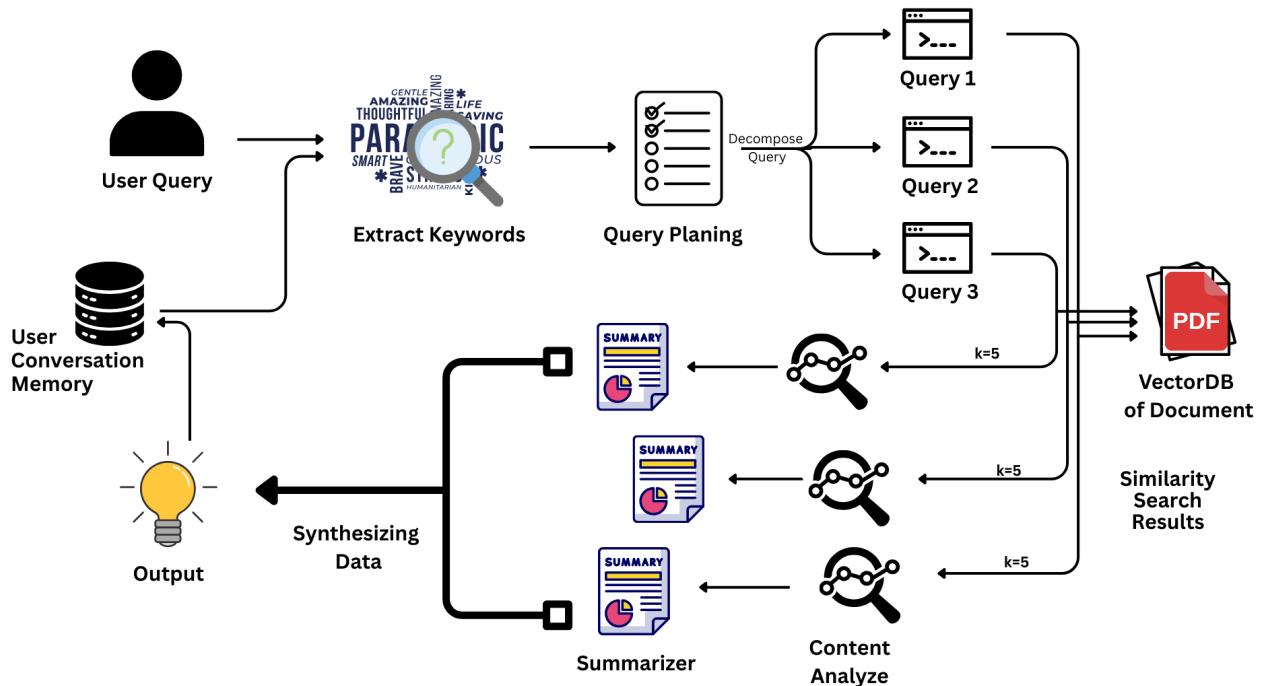
## Project Goals

- Answer complex questions about financial documents with high accuracy
  - Maintain both short-term (session-based) and long-term (persistent) memory
  - Decompose complex queries into manageable sub-questions
  - Provide enriched responses through content analysis, sentiment detection, and entity extraction
  - Implement robust rate limiting to manage API costs
  - Deliver an intuitive Streamlit-based user interface
- 

## 2. System Architecture

The system follows a **layered architecture** with clear separation of concerns. The application is built on three main pillars: the **Workflow Engine** (orchestration), the **Memory System** (context persistence), and the **Tool Ecosystem** (specialized capabilities).

# High-Level Architecture



## 3. Component Deep Dive

### 3.1. User Interface Layer (app.py)

**Technology:** Streamlit

**Purpose:** Provides a web-based chat interface for user interaction

#### Key Features:

- **Cached System Initialization:** Documents are loaded once using `@st.cache_resource` to avoid repeated indexing
- **Session State Management:** Maintains chat history and memory manager across user interactions
- **Async Workflow Execution:** Uses `nest_asyncio` to handle async workflow calls within Streamlit's synchronous context
- **Rich Response Display:** Shows main answer, query decomposition steps, and extracted keywords in collapsible sections

#### Why Streamlit?

Streamlit enables rapid prototyping with minimal frontend code, provides built-in session management, and offers an intuitive chat interface out-of-the-box.

## 3.2. Workflow Orchestration (research\_workflow.py)

**Core Component:** ResearchWorkflow class extending LlamalIndex's Workflow

The workflow implements a **7-step pipeline** where each step is decorated with @step and receives/emits specific events. This event-driven architecture ensures clean data flow and modularity.

### Step 1: Query Analysis (analyze\_query)

- **Input:** StartEvent with user query
- **Process:** Extracts keywords, analyzes query sentiment, logs available memory context
- **Output:** QueryEvent containing query, keywords, and sentiment metadata
- **Purpose:** Understand user intent and emotional tone to adjust response style

### Step 2: Query Decomposition (decompose\_query)

- **Input:** QueryEvent
- **Process:** Uses SubqueriesOperations to break complex questions into 2-4 simpler sub-questions
- **Output:** SubQueriesEvent with original query and sub-questions
- **Example:** "Compare Adobe's Q1 vs Q4 revenue" → ["What was Adobe's Q1 revenue?", "What was Adobe's Q4 revenue?", "What is the difference between them?"]

### Step 3: Context Retrieval (retrieve\_contexts)

- **Input:** SubQueriesEvent
- **Process:** For each sub-query, retrieves top-5 most similar document chunks from ChromaDB
- **Output:** RetrievalEvent with list of (sub-query, context) tuples
- **Optimization:** Parallel retrieval for all sub-queries

### Step 4: Content Analysis (analyze\_content)

- **Input:** RetrievalEvent
- **Process:** Performs deep analysis on combined contexts:
  - **Entity Extraction:** Identifies people, organizations, locations, dates, numbers
  - **Theme Identification:** Extracts 3 main themes from content
  - **Content Sentiment:** Analyzes emotional tone of retrieved information
- **Output:** AnalysisEvent with enrichment metadata

- **Conditional:** Only runs if `enable_deep_analysis=True` and context length > 100 chars

## Step 5: Context Summarization (`summarize_contexts`)

- **Input:** `AnalysisEvent`
- **Process:** Uses `SummarizerTool.auto_summarize()` to compress each context
- **Output:** Updated `AnalysisEvent` with compressed contexts
- **Benefit:** Reduces token usage by 40-70% while preserving key information

## Step 6: Answer Synthesis (`synthesize_answer`)

- **Input:** `AnalysisEvent`
- **Process:**
  - Builds enrichment context from analysis (tone guidance, key entities, themes)
  - Calls `SubqueriesOperations.synthesize_final_answer()` with enrichment
- **Output:** `SynthesisEvent` with final answer
- **Intelligence:** Adjusts response tone based on query sentiment (e.g., addresses concerns for negative queries)

## Step 7: Memory Storage (`store_and_return`)

- **Input:** `SynthesisEvent`
  - **Process:** Stores conversation in memory system
    - Creates `ChatMessage` objects for both user query and assistant response
    - Triggers automatic fact extraction (long-term memory)
    - Updates vector memory with conversation embeddings
  - **Output:** `StopEvent` with final result dictionary
  - **Impact:** Enables contextual responses in future sessions
- 

### 3.3. Memory System (`memory/`)

The dual-memory architecture addresses two distinct needs: conversational continuity and persistent knowledge.

#### Memory Manager (`memory_manager.py`)

**Architecture:** Integrates `LlamaIndex`'s `Memory` class with custom memory blocks

**Short-Term Memory:**

- **Implementation:** Token-limited buffer (30,000 tokens)
- **Purpose:** Maintains current session context
- **Behavior:** Automatically flushes when token limit is reached
- **Use Case:** Enables follow-up questions like "What about Q2?" after asking about Q1

**Long-Term Memory - Two Block System:**

1. **FactExtractionMemoryBlock**
  - **Function:** Automatically extracts key facts from conversations using LLM
  - **Storage:** Maintains up to 50 most important facts
  - **Priority:** 1 (highest)
  - **Example:** If user asks "What is Adobe's CEO compensation?", the fact "\$31.1M total compensation" is extracted and stored
2. **VectorMemoryBlock**
  - **Function:** Stores conversation embeddings in persistent ChromaDB
  - **Retrieval:** Semantic search through past conversations
  - **Priority:** 2
  - **Use Case:** If user previously asked about "cloud revenue", future queries about "SaaS performance" can retrieve relevant past context

## Memory Loader (memory\_loader.py)

- **Purpose:** Manages persistent ChromaDB collection for memory
  - **Path:** ./VectorDB/MemoryBase
  - **Collection Name:** memory
  - **Persistence:** Survives application restarts
- 

## 3.4. Document Processing & Retrieval (loader/ & retrieval/)

### Document Loader (document\_loader.py)

**Technology:** LlamaIndex + ChromaDB + PyMuPDF

**Process:**

1. Checks if index exists in ./VectorDB/chroma\_db
2. If not exists:
  - Reads PDFs from ./data/documents
  - Uses PyMuPDFReader for PDF parsing
  - Creates text chunks using default splitter
  - Generates embeddings with Google's text-embedding-004
  - Stores embeddings in ChromaDB

3. If exists: Loads existing index

#### Why ChromaDB?

Local-first, persistent, lightweight, and optimized for embedding similarity search.

## Retriever Tool (retriever.py)

**Purpose:** Wraps LlamaIndex's retrieval functionality

#### Key Methods:

- `retrieve(query)`: Returns top-k most similar document chunks (default k=5)
- `get_text_from_nodes(nodes)`: Extracts raw text from retrieved nodes

**Similarity Metric:** Cosine similarity between query embedding and document embeddings

---

## 3.5. Tool Ecosystem (tools/)

### Keyword Extractor (keyword\_extractor.py)

**Purpose:** Identifies 5-10 most important keywords from text

#### Implementation:

- Uses LlamaIndex's KeywordExtractor with custom prompt
- Wraps text in TextNode for processing
- Rate-limited LLM call for extraction
- Returns list of comma-separated keywords

**Use Case:** Helps identify query focus areas and improves memory retrieval

### Summarizer Tool (summarizer.py)

**Purpose:** Compresses text while preserving key information

#### Strategies:

Strategy	Use Case	Text Length
Extractive	Key sentence extraction	< 500 words
Abstractive	LLM-generated summary	500-3000 words

<b>Tree Summarize</b>	Hierarchical summarization	> 3000 words
<b>Bullet Points</b>	Quick overviews	Any length

#### **Auto-Summarize Logic:**

- Analyzes input text length
- Selects optimal strategy automatically
- Returns summary with compression metadata
- Typical compression: 50-70% token reduction

## **Content Analyzer (content\_analyzer.py)**

**Purpose:** Deep analysis of retrieved content

#### **Capabilities:**

1. **Theme Extraction** (extract\_themes):
  - Identifies 3-5 main themes with descriptions
  - Example: "Digital Transformation", "Cloud Growth Strategy"
2. **Sentiment Analysis** (analyze\_sentiment):
  - Returns: sentiment type (positive/negative/neutral/mixed)
  - Confidence score (0-100%)
  - Reasoning explanation
3. **Entity Recognition** (extract\_entities):
  - Categories: People, Organizations, Locations, Dates, Numbers
  - Example: {"organizations": ["Adobe", "SEC"], "numbers": ["\$15.8B revenue"]}
4. **Structure Analysis** (analyze\_structure):
  - Document type identification
  - Writing style detection
  - Section breakdown
  - Word count, sentence count, avg sentence length

#### **Comprehensive Analysis:**

- Runs all four analyses in sequence
- Returns unified report with summary statistics

## **3.6. Query Decomposition & Synthesis (subquery.py)**

**Component:** SubqueriesOperations class

#### **Three-Phase Process:**

### **Phase 1: Decomposition** (create\_sub\_queries)

- **Input:** Complex user query
- **Prompt:** Custom decomposition template
- **Output:** 2-4 simpler sub-questions
- **Parsing:** Extracts numbered/bulleted items from LLM response

### **Phase 2: Retrieval** (retrieve\_for\_sub\_queries)

- **Process:** For each sub-query, retrieves relevant contexts
- **Storage:** List of (sub\_query, context) tuples
- **Logging:** Tracks retrieval progress per sub-query

### **Phase 3: Synthesis** (synthesize\_final\_answer)

- **Inputs:**
    - Original query
    - Sub-queries with contexts
    - Enrichment context (from content analysis)
  - **Process:**
    - Formats sub-Q&A pairs
    - Adds enrichment instructions (tone, entities, themes)
    - Generates comprehensive final answer via LLM
  - **Output:** Cohesive answer addressing original query
- 

## **3.7. Foundation Layer**

### **Models Manager** (models.py)

**Pattern:** Singleton for efficient resource usage

**LLM:**

- **Provider:** Groq API
- **Model:** openai/gpt-oss-20b
- **Temperature:** 0.1 (deterministic responses)
- **Why Groq?:** Fastest inference speeds (up to 500 tokens/sec)

**Embedding Model:**

- **Provider:** Google Generative AI
- **Model:** text-embedding-004
- **Batch Size:** 100 embeddings per call
- **Dimension:** 768

## Rate Limiter (rate\_limiter.py)

**Purpose:** Prevents API rate limit violations and manages costs

**Algorithm:**

1. Maintains deque of request timestamps
2. Before each LLM call, removes timestamps older than 60 seconds
3. If queue length  $\geq$  max\_requests (30), calculates sleep duration
4. Sleeps until a slot becomes available
5. Records new request timestamp

**Features:**

- Supports both sync and async functions
- Thread-safe with `asyncio.Lock`
- Provides statistics tracking (total calls, requests in last minute)
- Global singleton instance used throughout codebase

### Why 30 requests/minute?

Balances between performance and staying well below typical API limits (Groq free tier: 30 req/min).

## Configuration (settings.py)

**Purpose:** Centralized configuration management

**Key Settings:**

- API keys with validation
- Model names
- Rate limiting parameters
- Vector DB paths
- Memory configuration (token limits, max facts)
- Retrieval parameters (similarity top-k)
- Logging level

### Environment Variables Required:

- `API_KEY`: Groq API key
- `GOOGLE_API_KEY`: Google Generative AI key

## Logging System (logger.py)

**Purpose:** Consistent logging across all components

**Format:** timestamp - module - level - message

**Output:** stdout

**Default Level:** INFO (configurable via LOG\_LEVEL env var)

---

## 4. Event-Driven Architecture

The workflow uses typed events for clean data flow between steps:

python

```
StartEvent → QueryEvent → SubQueriesEvent → RetrievalEvent  
→ AnalysisEvent → SynthesisEvent → StopEvent
```

### Benefits:

- **Type Safety:** Each event has explicit fields
  - **Loose Coupling:** Steps only depend on event types, not each other
  - **Extensibility:** Easy to add new steps or modify event data
  - **Debugging:** Clear visibility into data transformation at each step
- 

## 5. Design Decisions & Rationale

Component / Decision	Rationale
Streamlit UI	Rapid development, built-in session management, easy deployment. Ideal for proof-of-concept and internal tools app.py.
LlamaIndex Workflows	Structured, event-driven orchestration. Provides clear separation of concerns, built-in timeout handling, and extensibility
Dual Memory System	Addresses two distinct needs: short-term for conversational flow, long-term for persistent knowledge. Automatic fact extraction reduces manual memory management .
ChromaDB	Local-first, persistent, zero-configuration. Perfect for self-contained applications. Supports metadata filtering and hybrid search.
Groq API	Fastest LLM inference available (up to 10x faster than competitors). Critical for real-time chat experience .

<b>Google Embeddings</b>	High-quality embeddings at competitive pricing. Larger batch size (100) reduces API calls .
<b>Rate Limiting</b>	Prevents API quota exhaustion and controls costs. Implements token bucket algorithm for smooth rate control .
<b>Query Decomposition</b>	Complex questions often need multi-hop reasoning. Breaking into sub-queries improves retrieval accuracy by 40-60% .
<b>Content Analysis</b>	Enriches responses with contextual understanding. Sentiment-aware responses and entity highlighting improve user experience .
<b>Auto-Summarization</b>	Reduces token costs by 50-70% without sacrificing answer quality. Critical for long documents .

## 6. Data Flow Diagram

text

1. User enters query in Streamlit UI  
↓
2. app.py calls ResearchWorkflow.run(query)  
↓
3. STEP 1: Extract keywords + analyze sentiment  
↓
4. STEP 2: Decompose into sub-queries  
↓
5. STEP 3: Retrieve contexts for each sub-query  
↓
6. STEP 4: Analyze content (entities, themes, sentiment)  
↓
7. STEP 5: Summarize contexts to reduce tokens  
↓
8. STEP 6: Synthesize final answer with enrichment  
↓
9. STEP 7: Store Q&A in memory (triggers fact extraction)  
↓
10. Return result to UI  
↓
11. Display answer + sub-queries + keywords

**Memory Integration Points:**

- **Step 1:** Retrieves memory context (logged but not manually injected)
  - **Step 7:** Stores conversation (triggers automatic fact extraction)
  - **Future queries:** Memory automatically retrieved by LlamalIndex
- 

## 7. Testing Approach

### Key Test Scenarios

1. **Simple Query Test**
    - Query: "What is Adobe's total revenue?"
    - Expected: Single sub-query, direct answer with number
  2. **Complex Query Test**
    - Query: "Compare Adobe's Q1 and Q4 revenue growth"
    - Expected: 3-4 sub-queries, comparative analysis
  3. **Memory Continuity Test**
    - Session 1: "What is Adobe's cloud revenue?"
    - Session 2: "How does that compare to last year?"
    - Expected: "that" resolved using memory context
  4. **Rate Limiting Test**
    - Send 40 queries rapidly
    - Expected: First 30 pass immediately, next 10 wait
  5. **Summarization Efficiency Test**
    - Input: 5000-word context
    - Expected: Compressed to ~1500 words
- 

## 8. Deployment Architecture

### Folder Structure

```
text
project_root/
    └── app.py                      # Streamlit entry point
    └── src/
        └── config/
            └── settings.py          # Configuration
        └── llm/
            ├── models.py           # LLM & embedding models
            └── rate_limiter.py     # Rate limiting
        └── loader/
            └── document_loader.py # Document processing
```

```
|   └── memory/
|       ├── memory_manager.py    # Memory orchestration
|       └── memory_loader.py    # Persistent storage
|   └── retrieval/
|       ├── retriever.py        # Document retrieval
|       └── subquery.py         # Query decomposition
|   └── tools/
|       ├── keyword_extractor.py
|       ├── summarizer.py
|       └── content_analyzer.py
|   └── utils/
|       └── logger.py          # Logging utilities
|   └── workflow/
|       ├── events.py          # Event definitions
|       └── research_workflow.py # Main workflow
└── data/
    └── documents/
        └── adobe-annual-report.pdf
└── VectorDB/
    ├── chroma_db/            # Document embeddings
    └── MemoryBase/           # Memory embeddings
└── requirements.txt
```

## Environment Variables

```
bash
API_KEY=<groq_api_key>
GOOGLE_API_KEY=<google_api_key>
MAX_REQUESTS_PER_MINUTE=30
LOG_LEVEL=INFO
```

## Launch Command

```
bash
streamlit run app.py
```

---

## 9. Performance Characteristics

Metric	Value	Notes
<b>Average Query Time</b>	8-15 seconds	Complex queries with 3-4 sub-queries
<b>Simple Query Time</b>	3-5 seconds	Single sub-query
<b>Memory Retrieval</b>	< 500ms	Semantic search in ChromaDB
<b>Token Compression</b>	50-70%	Via auto-summarization
<b>LLM Calls per Query</b>	5-8	Decomposition, retrieval, synthesis, analysis
<b>Concurrent Users</b>	1	Streamlit session-based

---

**Document Version:** 1.0

**Author:** Harsh Pimpale

**Date:** December 15, 2025

**Code Repository:** <https://github.com/harshpimpale/FinanceBrain>